



AI spend governance audit

Every AI call must justify its cost.

Audit, simulate, enforce. Control spend before tokens become budget burn.

SYNTHETIC 500K WAL

RUST REPLAYED

CALYBRIS SHADOW

Verdict

ACTIONABLE COST LEAKAGE

Client	Synthetic Enterprise Replay
Run	9325b3257c2f3f91
Period	500k deterministic replay

Executive snapshot

Synthetic enterprise WAL replay. Use as a format and capability demo, not a customer savings claim. 500,000 synthetic WAL rows

Evidence maturity

89/100 Shadow-ready

Readiness inputs

Calls	500,000
Coverage	100.00%
Verdict	Cost leakage found

pilot gate target: 500+ replayed calls

TOTAL CALLS

500,000

Audited decision rows

CURRENT SPEND

\$3,196.55

Policy-selected model cost

AVOIDABLE SPEND

\$1,599.97

Requested minus selected

SAVINGS RATE

33.36%

Catalog estimate, not a claim

AT 100K CALLS/MO

\$319.99

Linear demo projection


AT 1M CALLS/MO

\$3,199.94


annualized \$38,399.26

Spend shape

Requested baseline versus GOVERIS policy and a balanced replay.



Policy decision mix



What-if policy simulation

Projection layer. Treat as pilot planning until replayed against real logs with outcome checks.

Scenario replay

				savings rate	100k-call range
Observed policy	Observed audit re...		33.36%	\$256.00-\$383.99	shadow-ready
Conservative	Low		20.00%	\$153.49-\$230.23	shadow-ready
Balanced	Medium		33.36%	\$256.00-\$383.99	shadow-ready
Aggressive	High		38.00%	\$291.63-\$437.44	shadow-ready





Recommended next actions

- Review blocked call 675863fbd2 (enterprise-acme/support-agent): net value -29.5000 cents, policy risk or budget guardrail triggered a block decision; confirm product-owner rules before enforcement.
- Review blocked call f458b0ebe8 (internal-tools/evaluation-pipeline): net value -22.8000 cents, policy risk or budget guardrail triggered a block decision; confirm product-owner rules before enforcement.
- Inspect call 0afa88a861 (enterprise-acme/coding-agent): gpt-4o -> gpt-4o-mini saved cost with confidence 0.83; verify answer quality before enabling automatic downgrade.
- This demo summary was computed by replaying the full 500,000-row synthetic WAL through the Rust audit command.
- Review downgraded calls first: GOVERIS found 116290 requests where a cheaper model satisfied the quality floor.

Attribution and audit trail

Where spend was created, where GOVERIS found avoidable cost, and what evidence was locked.

Priority calls to inspect				
Call	Slice	Decision	Route	Why it matters
675863fbd2	enterprise-acme support-agent	BLOCK	gpt-4o -> blocked	Policy risk or budget guardrail triggered a block decision conf 0.44; net -29.5000 cents; context hash phash_refund_ab...
f458b0ebe8	internal-tools evaluation-pipeline	BLOCK	gpt-4o -> blocked	Policy risk or budget guardrail triggered a block decision conf 0.51; net -22.8000 cents; context hash phash_eval_bulk...
0afa88a861	enterprise-acme coding-agent	DOWN GRADE	gpt-4o -> gpt-4o-mini	Cheaper model satisfied the configured quality floor conf 0.83; net 36.0520 cents; context hash phash_code_medi...
ba102e76ee	northstar-ai sales-assistant	DOWN GRADE	gpt-4o -> gpt-4o-mini	Cheaper model satisfied the configured quality floor conf 0.86; net 7.2190 cents; context hash phash_sales_sho...
15a23dbd57	enterprise-acme support-agent	DOWN GRADE	gpt-4o -> gpt-4o-mini	Cheaper model satisfied the configured quality floor conf 0.91; net 11.3240 cents; context hash phash_support_l...



Methodology notes	Interpretation guardrail
- CALYBRIS treats each LLM call as an investment decision: expected value minus model cost minus risk pressure. - Decision score is reported as net value = expected_value_cents - estimated_cost_cents - risk_penalty_cents. - Savings compare requested-model spend against the policy-selected model using the CALYBRIS model catalog. - What-if rows are scenario simulations, not realized production savings. - Projection ranges widen when the replay sample is small or audit-trail coverage is incomplete. - Quality preservation still requires outcome labels or human review samples. - This is a synthetic enterprise demo: useful for reviewing report structure and GOVERIS/CALYBRIS capabilities, not for customer savings proof.	GOVERIS estimates cost leakage from the model catalog and decision log. What-if savings are planning projections until replayed on real logs with outcome checks.

Evidence lock

Reproducible artifact trail for exactly this audit run.

Run ID

9325b3257c2f3f91

Generated

2026-06-17 17:09 UTC

Input SHA-256

0f9cc312208a48bfd8159adca99b81dc5794b782f3f77654684bf08bdc0573e

Summary SHA-256

165216c8afe389be70d74f9f63b94aa551efbc96fc2c9699ea7938e707d3fb92

What this proves

- same input file can be replayed
- summary artifact is locked
- not a production savings claim

Evidence lock answers: which rows were audited, when, and with which summary artifact.

30-day pilot proof plan

How this audit turns into enforceable spend governance.

01 Collect

500-5,000 anonymized calls with tenant, use case, tokens, risk, and confidence.

02 Replay

Run conservative, balanced, and aggressive what-if policies offline.

03 Review

Label high-risk, blocked, downgraded, and low-confidence decisions.

04 Enforce

Ship auth, rate limits, budgets, rollback, and audit retention.

Paid pilot intake checklist

Client sends	500-5,000 anonymized JSON/CSV decision rows. API keys, secrets, and raw PII are not required.	Minimum fields: tenant_id, use_case, requested_model, selected_model, input_tokens, output_tokens, estimated_cost, action, confidence.
Optional context	Prompt hash, prompt category, human-review label, latency, success/failure flag, and user-facing quality score.	If prompt text is sensitive, send a hash or short redacted preview. GOVERIS keeps missing context explicit.
GOVERIS returns	PDF audit, JSON summary, call review queue, tenant/use-case attribution, what-if policy replay, and next-step recommendations.	Savings remain planning estimates until replayed on larger logs with outcome labels or human review.
Commercial scope	Starter: sample audit. Standard: deeper replay and charts. Advanced: larger log review plus handoff notes.	Source code means client-specific helper script or notebook only when explicitly included; GOVERIS proprietary code is not transferred.

How to start the audit safely

Upload anonymized logs in the Upwork workroom. GOVERIS does not need provider API keys.

01 Share schema

JSON/CSV fields and model catalog notes.

02 Run replay

GOVERIS produces policy scenarios offline.

03 Review queue

Inspect blocked, downgraded, and costly

04 Deliver pack

PDF audit, JSON summary, and next actions.

Pilot acceptance gates

Enforcement remains locked until evidence, quality, and rollback controls are reviewable.

01 EVIDENCE

>=99% replay coverage

Input and summary hashes retained.

02 DATA QUALITY

Schema exceptions logged

No silent defaults or unknown models.

03 QUALITY

Risk queue reviewed

Downgrades checked against outcomes.

04 CONTROL

Rollback verified

Budgets, auth, and limits exercised.

GOVERIS

Govern. Verify. Operate.

Evidence-first AI spend governance, powered by the CALYBRIS decision engine.

01

AUDIT

Make model spend attributable.

02

SIMULATE

Replay policy before enforcement.

03

ENFORCE

Ship only reviewed controls.